

# Tree-based Methods

## The Data Set

```
library(tree)
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(gbm)
```

```
## Loaded gbm 2.1.8.1
```

```
biodeg = read.csv("biodeg.csv", stringsAsFactors=TRUE)
head(biodeg)
```

```
##      SpMaxL   JDze nHM F01NN F04CN NssssC nCb      C nCp nO F03CN   SdssC HyWiBm
## 1  3.919 2.6909   0     0     0       0  0 31.4   2  0     0  0.000  3.106
## 2  4.170 2.1144   0     0     0       0  0 30.8   1  1     0  0.000  2.461
## 3  3.932 3.2512   0     0     0       0  0 26.7   2  4     0  0.000  3.279
## 4  3.000 2.7098   0     0     0       0  0 20.0   0  2     0  0.000  2.100
## 5  4.236 3.3944   0     0     0       0  0 29.4   2  4     0 -0.271  3.449
## 6  4.236 3.4286   0     0     0       0  0 28.6   2  4     0 -0.275  3.313
##      LOC   SM6L F03CO     Me     Mi nNN nArNO2 nCRX3 SpPosABp nCIR B01CBr B03CCL
## 1 2.550 9.002     0 0.960 1.142   0       0     0     1.201   0     0     0
## 2 1.393 8.723     1 0.989 1.144   0       0     0     1.104   1     0     0
## 3 2.585 9.110     0 1.009 1.152   0       0     0     1.092   0     0     0
## 4 0.918 6.594     0 1.108 1.167   0       0     0     1.024   0     0     0
## 5 2.753 9.528     2 1.004 1.147   0       0     0     1.137   0     0     0
## 6 2.522 9.383     1 1.014 1.149   0       0     0     1.119   0     0     0
##      N073 SpMaxA Psiild B04CBr      SdO  TI2L nCrt C026 F02CN nHDon SpMaxBm PsiiA nN
## 1    0  1.932  0.011     0  0.000 4.489   0    0     0     0     2.949 1.591  0
## 2    0  2.214 -0.204     0  0.000 1.542   0    0     0     0     3.315 1.967  0
## 3    0  1.942 -0.008     0  0.000 4.891   0    0     0     1     3.076 2.417  0
## 4    0  1.414  1.073     0  8.361 1.333   0    0     0     1     3.046 5.000  0
## 5    0  1.985 -0.002     0 10.348 5.588   0    0     0     0     3.351 2.405  0
## 6    0  1.980 -0.008     0 10.276 4.746   0    0     0     0     3.351 2.556  0
##      SM6Bm nArCOOR nX class
## 1 7.253     0  0    RB
## 2 7.257     0  0    RB
## 3 7.601     0  0    RB
## 4 6.690     0  0    RB
## 5 8.003     0  0    RB
## 6 7.904     0  0    RB
```

# Tasks

## Training and Test Data


1. Randomly divide the data set into two halves, and save them in two data frames named train and test.

```
set.seed(769)
index <- sample(rep(1:2, length.out=nrow(biodeg)))
train <- biodeg[index ==1, ]
test <- biodeg[index ==2, ]
```

## Classification trees

2. Fit an unpruned classification tree to the training data. Plot it (as pretty as you can). Identify three most important variables from this classification tree.

```
r = tree(class ~ ., data=train)
plot(r)
text(r, pretty=1, cex=0.5)
```



For the heights of splits are proportional to deviance reductions, so we can see from the tree that the 3 most important variables are F02CN, F04CN and nCb.

3. Compute the training and test errors. Write a function `errors(fit, fhat.tree, train, test)` where `fit` is the output of `tree()` and `fhat.tree` is a function that uses `fit` and computes the class labels for a data set (an argument of `fhat.tree`).

```
fhat.tree <- function(fit, dataset){
  p = predict(fit, dataset)
  yhat = levels(dataset$class)[apply(p, 1, which.max)]
  yhat
}


errors <- function(fit, fhat.tree, train, test){
  train.error = mean(fhat.tree(fit, train) != train$class)
  test.error = mean(fhat.tree(fit, test) != test$class)
  c(train.error, test.error)
}

(errors(r, fhat.tree, train, test))
```

```
## [1] 0.1098485 0.1726755
```

4. Consider pruning the tree using cross-validation with deviance. Produce a pruned tree based by selecting a cost-complexity parameter value, and plot it. Compute the training and test errors for this pruned tree. Do you think the pruning helps?

```
cv.r = cv.tree(r)
j.min = which.min(cv.r$dev)
k = cv.r$k[j.min]
r2 = prune.tree(r, k=k)
plot(r2)
text(r2, pretty=0)
```




```
(errors(r2, fhat.tree, train, test))
```

```
## [1] 0.1856061 0.2314991
```

The tree is considerably smaller, which is computationally efficient, so I think pruning helps in the trade-off of accuracy.

5. Consider pruning the tree using cross-validation with misclassification rates. Produce a pruned tree by selecting a tree size, and plot it. Compute the training and test errors for this pruned tree. Do you think the pruning helps?

```
cv.r = cv.tree(r, FUN = prune.misclass)
j.min = which.min(cv.r$dev)
size = cv.r$size[j.min]
r3 = prune.tree(r, best=size)
plot(r3)
text(r3, pretty=1, cex=0.5)
```



```
(errors(r3, fhat.tree, train, test))
```

```
## [1] 0.1098485 0.1726755
```

It produced the exactly the same tree as the unpruned tree, so I don't think pruning helps this time.

## Bagging

6. Produce a Bagging model for the training data with 500 trees constructed. What are the three most important variables, in terms of decreasing the Gini index, according to Bagging?

```
p = ncol(biodeg) - 1
(r = randomForest(class ~ ., data=train, mtry=p, importance=TRUE, ntree=500))
```

```
##
## Call:
## randomForest(formula = class ~ ., data = train, mtry = p, importance = TRUE,
ntree = 500)
##
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 41
##
##           OOB estimate of  error rate: 15.72%
## Confusion matrix:
##      NRB  RB class.error
## NRB 325  33  0.09217877
## RB   50 120  0.29411765
```

```
round(importance(r), 2)
```

##	NRB	RB	MeanDecreaseAccuracy	MeanDecreaseGini
## SpMaxL	5.02	24.49	24.04	20.27
## JDze	7.24	6.21	9.68	5.79
## nHM	0.56	9.50	9.99	1.39
## F01NN	2.38	2.65	2.95	0.26
## F04CN	1.78	9.96	9.10	2.72
## NssssC	4.64	7.37	7.21	1.38
## nCb	3.63	2.69	5.14	0.80
## C	1.52	2.96	3.52	5.35
## nCp	1.68	11.36	9.62	2.91
## nO	5.69	5.03	8.09	3.37
## F03CN	-1.46	10.69	9.43	2.46
## SdssC	15.45	17.22	21.81	12.51
## HyWiBm	13.18	6.66	15.44	7.23
## LOC	6.52	8.14	10.77	5.72
## SM6L	9.05	4.81	11.23	4.52
## F03CO	8.40	8.55	11.61	4.87
## Me	8.90	2.90	9.93	5.56
## Mi	9.03	8.44	12.24	7.68
## nNN	2.46	1.42	2.31	0.19
## nArNO2	2.73	2.24	3.16	0.15
## nCRX3	0.00	0.00	0.00	0.01
## SpPosABp	13.03	9.53	16.79	11.43
## nCIR	0.94	3.24	3.41	0.48
## B01CBr	1.00	0.00	1.00	0.03
## B03CCl	0.38	4.54	3.66	0.36
## N073	1.42	-2.46	0.49	0.28
## SpMaxA	10.01	12.34	16.80	12.35
## Psiild	12.38	-3.99	10.91	3.72
## B04CBr	-1.42	0.00	-1.42	0.05
## SdO	6.12	4.66	7.63	3.70
## TI2L	8.67	4.58	9.87	5.44
## nCrt	4.96	8.94	9.85	1.60
## C026	3.73	4.17	5.52	1.09
## F02CN	8.06	22.24	21.84	13.05
## nHDon	2.31	4.44	5.33	2.36
## SpMaxBm	21.34	20.33	31.57	42.09
## PsiiA	12.14	6.85	14.31	8.81
## nN	6.48	14.32	13.68	6.51
## SM6Bm	20.58	10.95	24.09	17.33
## nArCOOR	2.00	8.99	8.15	3.08
## nX	-1.75	11.42	10.72	1.60

SpMaxBm, SM6Bm and SpMaxL are the three most important variables according to bagging.

7. Compute both the training and test errors of this Bagging predictor. Is your test error similar to the OOB estimate? Do you think Bagging helps prediction here?

```
yhat = predict(r, train)
(mean(train$class != yhat))
```

```
## [1] 0
```

```
yhat = predict(r, test)
(mean(test$class != yhat))
```

```
## [1] 0.1442125
```

The test set error is pretty close to OOB error, and the error rate slightly decreased. So I think bagging helps prediction.

## Random Forests

8. Produce a Random Forest model with 500 trees constructed. What are the three most important variables, in terms of accuracy, according to Random Forest?

```
(r = randomForest(class ~ ., data=train, mtry=10, importance=TRUE, ntree=500))
```

```
##
## Call:
## randomForest(formula = class ~ ., data = train, mtry = 10, importance = TRUE,
ntree = 500)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 10
##
## OOB estimate of  error rate: 15.53%
## Confusion matrix:
##      NRB  RB class.error
## NRB 324  34  0.09497207
## RB   48 122  0.28235294
```

```
round(importance(r), 2)
```

##	NRB	RB	MeanDecreaseAccuracy	MeanDecreaseGini
## SpMaxL	10.59	18.62	19.60	18.49
## JDze	6.86	5.47	9.22	6.84
## nHM	5.02	9.36	10.63	3.23
## F01NN	1.59	4.04	4.34	0.36
## F04CN	5.08	9.67	9.98	3.03
## NssssC	4.75	6.92	7.36	2.21
## nCb	5.50	4.25	7.10	2.75
## C	4.36	1.66	4.79	6.31
## nCp	5.14	8.59	9.59	3.34
## nO	6.61	10.21	12.47	5.23
## F03CN	5.00	10.36	11.27	4.12
## SdssC	11.44	12.80	15.95	9.31
## HyWiBm	12.73	8.56	14.79	9.89
## LOC	2.46	8.59	7.80	5.86
## SM6L	8.83	9.82	13.78	8.26
## F03CO	8.05	8.98	12.81	5.10
## Me	6.94	3.10	8.21	5.82
## Mi	7.65	6.63	10.64	7.01
## nNN	1.15	1.00	1.41	0.25
## nArNO2	2.57	2.13	2.91	0.15
## nCRX3	0.00	1.00	1.00	0.03
## SpPosABp	13.36	11.66	18.13	14.47
## nCIR	2.24	3.43	4.36	1.22
## B01CBr	-1.61	1.00	-1.13	0.04
## B03CCl	-0.54	3.07	1.97	0.26
## N073	1.56	-1.74	0.98	0.20
## SpMaxA	10.23	12.85	15.96	14.42
## Psiild	6.91	-2.28	5.88	4.26
## B04CBr	-2.55	0.00	-2.56	0.07
## SdO	7.39	10.34	12.57	5.41
## TI2L	6.73	7.31	10.06	6.22
## nCrt	3.84	4.28	5.97	1.27
## C026	5.57	5.49	6.72	2.03
## F02CN	9.14	15.27	16.97	8.99
## nHDon	3.50	4.06	5.77	2.76
## SpMaxBm	14.15	15.78	20.84	22.69
## PsiiA	8.52	5.95	10.98	7.65
## nN	9.62	16.12	16.81	8.65
## SM6Bm	15.15	13.59	21.59	17.56
## nArCOOR	2.49	6.73	6.70	2.28
## nX	3.90	8.80	9.35	1.73

SpMaxBm, SpMaxL and SM6Bm are the most three important variables according to random forrest.

9. Compute both the training and test errors of this Random Forest predictor. Is your test error similar to the OOB estimate? Do you think the tweak used by Random Forest helps prediction here?

```
yhat = predict(r, train)
(mean(train$class != yhat))
```

```
## [1] 0
```



```
yhat = predict(r, test)
(mean(test$class != yhat))
```

```
## [1] 0.1347249
```

The test set error is pretty close to OOB error, and the error rate was even better than that with bagging. So I think it helps prediction.


## Boosting

10. Produce a Boosting model, with 500 trees constructed. What are the three most important variables, according to Boosting?

```
train$class = as.integer(train$class)-1
test$class = as.integer(test$class)-1
(r = gbm(class ~ ., data=train, distribution="bernoulli", n.trees=500))
```

```
## gbm(formula = class ~ ., distribution = "bernoulli", data = train,
##      n.trees = 500)
## A gradient boosted model with bernoulli loss function.
## 500 iterations were performed.
## There were 41 predictors of which 31 had non-zero influence.
```

```
summary(r)
```



##	var	rel.inf
##	SpMaxBm	SpMaxBm 17.06033867
##	SdssC	SdssC 8.73791894
##	SpMaxL	SpMaxL 8.33174383
##	SpPosABp	SpPosABp 6.53530162
##	HyWiBm	HyWiBm 5.54586333
##	PsiiA	PsiiA 5.40705813
##	Mi	Mi 5.13265840
##	F02CN	F02CN 4.50074247
##	SM6Bm	SM6Bm 4.23748234
##	LOC	LOC 4.16786118
##	nN	nN 3.90635146
##	SpMaxA	SpMaxA 3.77574882
##	C	C 3.03901731
##	JDze	JDze 2.97354214
##	TI2L	TI2L 2.62600574
##	Psiild	Psiild 1.76390280
##	Me	Me 1.50505290
##	F03CO	F03CO 1.37128421
##	nO	nO 1.26297730
##	nHDon	nHDon 1.22289433
##	SdO	SdO 1.18392919
##	NssssC	NssssC 1.14402089
##	nArCOOR	nArCOOR 0.85722592
##	C026	C026 0.81399057
##	nCrt	nCrt 0.78544697
##	SM6L	SM6L 0.74975356
##	nX	nX 0.48314202
##	nCb	nCb 0.36209309
##	nCp	nCp 0.34350445
##	nHM	nHM 0.11477904
##	F03CN	F03CN 0.05836838
##	F01NN	F01NN 0.00000000
##	F04CN	F04CN 0.00000000
##	nNN	nNN 0.00000000
##	nArNO2	nArNO2 0.00000000
##	nCRX3	nCRX3 0.00000000
##	nCIR	nCIR 0.00000000
##	B01CBr	B01CBr 0.00000000
##	B03CCl	B03CCl 0.00000000
##	N073	N073 0.00000000
##	B04CBr	B04CBr 0.00000000

SpMaxBm, SdssC and SpMaxL are the three most important variables according to boosting.

11. Compute both the training and test errors of this Boosting predictor. Do you think Boosting helps prediction here?

```
yhat = (predict(r, train, type="response") > 0.5)
```

```
## Using 500 trees...
```

```
(mean(train$class != yhat))
```

```
## [1] 0.0530303
```

```
yhat = (predict(r, test, type="response") > 0.5)
```

```
## Using 500 trees...
```

```
(mean(test$class != yhat))
```

```
## [1] 0.1480076
```

It helps in regards to the pruned decision tree, but the test error could be too optimistic compared to OOB error from bagging and random forest.

12. Demonstrate that Boosting can overfit.

```
r = gbm(class ~ ., data=train, distribution="bernoulli", n.trees=200000, n.cores=8 )  
yhat = (predict(r, train, type="response") > 0.5)
```

```
## Using 200000 trees...
```

```
(mean(train$class == yhat))
```

```
## [1] 1
```

```
yhat = (predict(r, test, type="response") > 0.5)
```

```
## Using 200000 trees...
```

```
(mean(test$class == yhat))
```

```
## [1] 0.8045541
```

## Summary

In this lab, we played around QSAR biodegradation data set which uses 41 molecular descriptors (SpMax\_L, etc.) to predict class labels (ready biodegradable (RB) and not ready biodegradable (NRB)).

We used the tree family to do the classification. We went from the foundation - a single decision tree, to bagging which over-samples the data set with replacement so that different models can be combined to have a unbiased estimate, to random forest with the tweak to use only some of the variables to find the optimal cutoff points, to boosting that gradually makes to a better prediction.

We found that SpMaxBm is the common important variable recognised by all the models, and all the models have a fair prediction of the class labels.