

# A2 - How to do Machine Learning

September 19, 2022

```
[ ]: # import packages
import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from ReliefF import ReliefF
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from autorank import autorank
```

As the missing values scatter across the data set, so we can't simply delete them. And data itself covers a range of different values so we rule out imputation and instead fill in a global constant not to undermine the variability while remain validity of the data.

```
[ ]: # read data
X = pd.read_csv('data_A2.csv').to_numpy()
y = pd.read_csv('labels_A2.csv').to_numpy().flatten()
# Fill the missing value with -1
X = np.nan_to_num(X, copy=True, nan=-1)
X.shape
```

```
[ ]: (999, 100)
```

Majority of learning methods won't behave well in the data set due to curse of dimensionality (100 features!). Thus it is important that before a model is trained, a feature selection technique is applied.

We can 1. calculate the correlation between each feature and the class label 2. use a classifier to test out the different combinations of features (random forest will be a good option as it can randomly impute the variables in the OOB cases, and then compare them to those not in the OOB) 3. rely on a dedicated feature selection method (PCR or Relief)

In our case, we opt for the dedicated feature selection method ReliefF. This is because we don't want to assume conditional independence upon the class label between features (rule out method number one) and any classifier won't function properly with these many features, so no point in testing feature importance with an underperformed classifier in the first place (rule out method number two).

```
[ ]: # set a random state to keep a reproducible output
RANDOM_STATE = 1234
np.random.seed(RANDOM_STATE)
fs = ReliefF(n_neighbors=1, n_features_to_keep=10)
X = fs.fit_transform(X, y)
# X_clean = X_fill[:,np.argsort(fi)[:10]]
```

We use tree-based classifiers (decision stump, unpruned tree, pruned tree, random forest) to establish a baseline result before we play around the data set (additive noise, multiplicative noise, etc.)

```
[ ]: def TreebasedClassifiers(X, y):
    scores = pd.DataFrame()

    ds = DecisionTreeClassifier(max_depth=1, random_state= RANDOM_STATE)
    scores["Decistion Stump"] = cross_val_score(ds, X, y, cv=10)

    # we set the max-depth equal to the number of features, so it will grow to
    ↳ the maximum depth (unpruned)
    dt = DecisionTreeClassifier(max_depth=10, random_state= RANDOM_STATE)
    scores["Decistion Tree"] = cross_val_score(dt, X, y, cv=10)

    # we set minimal_cost_complexity_pruning close enough to 0, so it will
    ↳ differ from decision stumps, but slight larger than 0, so it will differ
    ↳ from an unpruned decision tree (pruned)
    pt = DecisionTreeClassifier(random_state=RANDOM_STATE, ccp_alpha=0.005)
    scores["Decistion Tree (Prunned)"] = cross_val_score(pt, X, y, cv=10)

    rf = RandomForestClassifier(max_depth=10, random_state= RANDOM_STATE)
    scores["Random Forest"] = cross_val_score(rf, X, y, cv=10)

    result = autorank(scores, alpha=0.05, verbose=False)
    print(result)

TreebasedClassifiers(X, y)
```

RankResult(rankdf=

	meanrank	mean	std	ci_lower	ci_upper \
Decistion Stump	3.10	0.558525	0.065985	0.526063	0.590987
Decistion Tree (Prunned)	2.85	0.576525	0.050196	0.544063	0.608987
Decistion Tree	2.70	0.575525	0.056072	0.543063	0.607987
Random Forest	1.35	0.633606	0.040060	0.601144	0.666068

	effect_size	magnitude
Decistion Stump	0.0	negligible
Decistion Tree (Prunned)	-0.307041	small
Decistion Tree	-0.277644	small
Random Forest	-1.375519	large

```

pvalue=0.00023423524796781616
cd=None
omnibus=anova
posthoc=tukeyhsd
all_normal=True
pvals_shapiro=[0.31625089049339294, 0.8351843953132629, 0.02381780743598938,
0.7750453352928162]
homoscedastic=True
pval_homogeneity=0.5368611191270642
homogeneity_test=bartlett
alpha=0.05
alpha_normality=0.0125
num_samples=10
posterior_matrix=
None
decision_matrix=
None
rope=None
rope_mode=None
effect_size=cohen_d)

```

The decision stump performed understandably the worst as it only threshold on one feature which leads to underfitting.

```

[ ]: # additive normal noise
noise = np.random.normal(0, 0.2, np.shape(X))
X_addictive_noice = X + np.multiply(noise, np.average(X, axis=0))
TreebasedClassifiers(X_addictive_noice, y)

```

```

RankResult(rankdf=

```

	meanrank	mean	std	ci_lower	ci_upper	\
Decistion Stump	2.95	0.548525	0.056658	0.517434	0.579616	
Decistion Tree (Prunned)	2.95	0.550505	0.053456	0.519414	0.581596	
Decistion Tree	2.90	0.550505	0.043612	0.519414	0.581596	
Random Forest	1.20	0.624626	0.051881	0.593535	0.655717	

```


```

	effect_size	magnitude
Decistion Stump	0.0	negligible
Decistion Tree (Prunned)	-0.035944	negligible
Decistion Tree	-0.03916	negligible
Random Forest	-1.400924	large

```

pvalue=3.436984291080978e-05
cd=None
omnibus=anova
posthoc=tukeyhsd
all_normal=True
pvals_shapiro=[0.6545984148979187, 0.23763686418533325, 0.04130295291543007,
0.27765145897865295]

```

```

homoscedastic=True
pval_homogeneity=0.8920393834889994
homogeneity_test=bartlett
alpha=0.05
alpha_normality=0.0125
num_samples=10
posterior_matrix=
None
decision_matrix=
None
rope=None
rope_mode=None
effect_size=cohen_d)

```

All the classifiers remain stable, as the noise can be learned.

```

[ ]: # multiplicative normal noise
noise = np.random.normal(0, 0.2, np.shape(X))
X_multiplicative_noise = np.multiply(X, noise)
TreebasedClassifiers(X_multiplicative_noise, y)

```

```
RankResult(rankdf=
```

	meanrank	median	mad	ci_lower	ci_upper	\
Decistion Stump	2.95	0.490	0.01	0.45	0.505051	
Random Forest	2.60	0.495	0.025	0.38	0.52	
Decistion Tree (Prunned)	2.50	0.500	0.0	0.43	0.505051	
Decistion Tree	1.95	0.505	0.015	0.42	0.535354	

	effect_size	magnitude
Decistion Stump	0.0	negligible
Random Forest	-0.17713	negligible
Decistion Tree (Prunned)	-0.953874	large
Decistion Tree	-0.793671	medium

```

pvalue=0.3088962111457305
cd=1.483221853685529
omnibus=friedman
posthoc=nemenyi
all_normal=False
pvals_shapiro=[0.01277504488825798, 0.05152953788638115, 3.401902404220891e-07,
0.04103650152683258]
homoscedastic=False
pval_homogeneity=0.03401818284279286
homogeneity_test=levane
alpha=0.05
alpha_normality=0.0125
num_samples=10
posterior_matrix=
None

```

```

decision_matrix=
None
rope=None
rope_mode=None
effect_size=akinshin_gamma)

```

All the classifiers are worse as the noise can't be corrected due to randomised scalar.

```

[ ]: ## class noise
mask = np.random.binomial(1, 0.05, y.shape[0])
y_class_noise = abs(y - mask)
TreebasedClassifiers(X, y_class_noise)

```

```

RankResult(rankdf=

```

	meanrank	mean	std	ci_lower	ci_upper	\
Decistion Stump	3.55	0.564566	0.041128	0.541299	0.587832	
Decistion Tree	2.80	0.577636	0.038182	0.55437	0.600903	
Decistion Tree (Prunned)	2.15	0.594616	0.023738	0.57135	0.617882	
Random Forest	1.50	0.631667	0.047512	0.6084	0.654933	

```


```

	effect_size	magnitude
Decistion Stump	0.0	negligible
Decistion Tree	-0.329383	small
Decistion Tree (Prunned)	-0.894942	large
Random Forest	-1.510099	large

```

pvalue=0.00010545055797056721
cd=None
omnibus=anova
posthoc=tukeyhsd
all_normal=True
pvals_shapiro=[0.5788647532463074, 0.20331569015979767, 0.3019199073314667,
0.868999183177948]
homoscedastic=True
pval_homogeneity=0.2673054416861556
homogeneity_test=bartlett
alpha=0.05
alpha_normality=0.0125
num_samples=10
posterior_matrix=
None
decision_matrix=
None
rope=None
rope_mode=None
effect_size=cohen_d)

```

# 1 Summary

The multiplicative noise will have a big impact on the performance, while the other kinds of noise can be coped well.